

data-cleaning

December 7, 2024

1 Importing the required modules

```
[1]: # modules used for data handling and
# manipulation
import numpy as np
import pandas as pd

# modules used for data visualization
import matplotlib.pyplot as plt
import seaborn as sns

# modules used for encoding and data splitting
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import TargetEncoder, OrdinalEncoder, StandardScaler

# ignore warnings
import warnings
warnings.filterwarnings("ignore")
```

2 Reading the data

```
[2]: flight_df = pd.read_csv("DelayData.csv")
```

3 Initial Data Exploration

```
[3]: flight_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1201664 entries, 0 to 1201663
Data columns (total 61 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   depdelay                             1201664 non-null int64
 1   arrdelay                             1198458 non-null float64
 2   scheduleddepartdatetime              1201664 non-null object
 3   origin                               1201664 non-null object
```

4	dest	1201664	non-null	object
5	uniquecarrier	1201664	non-null	object
6	marketshareorigin	1201664	non-null	float64
7	marketsharedest	1201664	non-null	float64
8	hhiorigin	1201664	non-null	float64
9	hhidest	1201664	non-null	float64
10	nonhubairportorigin	1201664	non-null	int64
11	smallhubairportorigin	1201664	non-null	int64
12	mediumhubairportorigin	1201664	non-null	int64
13	largehubairportorigin	1201664	non-null	int64
14	nonhubairportdest	1201664	non-null	int64
15	smallhubairportdest	1201664	non-null	int64
16	mediumhubairportdest	1201664	non-null	int64
17	largehubairportdest	1201664	non-null	int64
18	nonhubairlineorigin	1201664	non-null	int64
19	smallhubairlineorigin	1201664	non-null	int64
20	mediumhubairlineorigin	1201664	non-null	int64
21	largehubairlineorigin	1201664	non-null	int64
22	nonhubairlinedest	1201664	non-null	int64
23	smallhubairlinedest	1201664	non-null	int64
24	mediumhubairlinedest	1201664	non-null	int64
25	largehubairlinedest	1201664	non-null	int64
26	year	1201664	non-null	int64
27	month	1201664	non-null	int64
28	dayofmonth	1201664	non-null	int64
29	dayofweek	1201664	non-null	int64
30	scheduledhour	1201664	non-null	int64
31	originairportid	1201664	non-null	int64
32	destairportid	1201664	non-null	int64
33	tailnum	1201664	non-null	object
34	capacity	1201664	non-null	int64
35	loadfactor	1201664	non-null	float64
36	numflights	1201664	non-null	float64
37	origincityname	1201664	non-null	object
38	originstate	1201664	non-null	object
39	distance	1201664	non-null	int64
40	monopolyroute	1201664	non-null	int64
41	temperature	1201204	non-null	float64
42	temp_ninfty_n10	1201664	non-null	int64
43	temp_n10_0	1201664	non-null	int64
44	temp_0_10	1201664	non-null	int64
45	temp_10_20	1201664	non-null	int64
46	temp_20_30	1201664	non-null	int64
47	temp_30_40	1201664	non-null	int64
48	temp_40_infty	1201664	non-null	int64
49	windspeed	1201204	non-null	float64
50	windspeedsquare	1201204	non-null	float64
51	windgustdummy	1201664	non-null	int64

```

52 windgustspeed          1201204 non-null float64
53 raindummy             1201664 non-null int64
54 raintracedummy        1201664 non-null int64
55 snowdummy             1201664 non-null int64
56 snowtracedummy        1201664 non-null int64
57 originmetropop         1201664 non-null int64
58 originmetrogdppercapita 1201664 non-null float64
59 destmetropop           1201664 non-null int64
60 destmetrogdppercapita   1201664 non-null float64

```

dtypes: float64(13), int64(41), object(7)

memory usage: 559.2+ MB

- The dataset has 1201664 rows and 61 columns.
- The dataset is partially pre-processed. Few of the categorical variables are already one-hot encoded.
- However, the many of categorical variables that are already encoded are ordinal in nature and one-hot encoding is not appropriate for it.
- It also appears that the `windspeed` column is transformed into `windspeedsquare`. It must be checked which one is better and only one of them must be retained.
- The first five rows of the dataset are shown below.

```
[4]: flight_df.head()
```

```

[4]:  depdelay  arrdelay  scheduleddepartdatetime  origin  dest  uniquecarrier  \
0         0        -4.0    08-Jan-2004 15:25:00    ELP    SAT                WN
1        -4         11.0    22-Jan-2004 14:40:00    ATL    MSY                DL
2         3         12.0    29-Jan-2004 12:25:00    DFW    JFK                DL
3        -3         24.0    14-Jan-2004 15:55:00    SEA    EWR                CO
4         0         -8.0    14-Jan-2004 18:40:00    SLC    RNO                OO

      marketshareorigin  marketsharedest  hhiorigin  hhidest  ...  \
0          0.618467          0.407567    0.417090  0.226878  ...
1          0.500757          0.096321    0.319589  0.196657  ...
2          0.060898          0.131962    0.296126  0.214357  ...
3          0.040522          0.347744    0.234712  0.249377  ...
4          0.506899          0.176493    0.341763  0.277364  ...

      windgustdummy  windgustspeed  raindummy  raintracedummy  snowdummy  \
0                 0              0.0         0                0          0
1                 0              0.0         0                0          0
2                 0              0.0         0                0          0
3                 0              0.0         1                0          0
4                 0              0.0         0                0          0

      snowtracedummy  originmetropop  originmetrogdppercapita  destmetropop  \
0                   0             702433              27314.633      1843927
1                   0             4802300              49081.773      1314721
2                   0             5689982              50588.563      18747431

```

3	0	3163703	57755.547	18747431
4	0	1030597	45043.602	385049

	destmetrogdppercapita
0	35005.234
1	48848.234
2	57295.402
3	57295.402
4	49079.727

[5 rows x 61 columns]

4 Data Cleaning

4.1 Handling null values

```
[5]: def display_cols_wt_na(df):
      print(df.isna().sum().loc[lambda x : x>0].sort_values(ascending = False)*100/
      ↪len(df))
```

```
[6]: display_cols_wt_na(flight_df)
```

```
arrdelay      0.266797
temperature   0.038280
windspeed     0.038280
windspeedsquare 0.038280
windgustspeed 0.038280
dtype: float64
```

- The `arrdelay` has 0.26% null values. It is also the target column and hence the **rows** with null value for `arrdelay` must be dropped.
- After dropping those rows we will check if the other columns yet have null values.

```
[7]: flight_df = flight_df[flight_df['arrdelay'].notna()]
      display_cols_wt_na(flight_df)
```

```
temperature   0.038383
windspeed     0.038383
windspeedsquare 0.038383
windgustspeed 0.038383
dtype: float64
```

- All the columns have 0.03% null values and hence can be imputed.
- Now, we will impute the missing values in the `temperature`, `windspeed`, `windspeedsquare`, and `windgustspeed` with the **mean** of the columns.
- `np.nanmean()` is used to compute mean of columns with null values.

```
[8]: to_impute_cols = ["temperature", "windspeed", "windspeedsquare",
      ↪"windgustspeed"]
```

```

for col in to_impute_cols:
    col_mean = np.nanmean(flight_df[col])
    flight_df[col] = flight_df[col].fillna(col_mean)

```

```
[9]: display_cols_wt_na(flight_df)
```

```
Series([], dtype: float64)
```

All null values are handled.

4.2 Dropping the columns

- There are few redundant columns and they need dropped.
- The `scheduleddepartdatetime` column stores the timestamp for each flight. However, the dataset also has all the individual components in separate columns and hence this column is dropped.
- The `originairportid` and `destairportid` columns have same information as the `origin` and `dest` columns. Hence, they are dropped.
- The `originstate` and `origincityname` columns have same information as `origin` column. While, this data is unknown for `dest`. Hence, those two columns are dropped.

Note: the time-series nature of data is not considered for the modelling.

```

[10]: cols_to_be_dropped = ["scheduleddepartdatetime",
                             "originairportid",
                             "destairportid",
                             "originstate",
                             "origincityname"]

flight_df = flight_df.drop(columns = cols_to_be_dropped,
                           axis = 1)

```

```
[11]: flight_df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Index: 1198458 entries, 0 to 1201663
Data columns (total 56 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   depdelay                             1198458 non-null  int64
1   arrdelay                             1198458 non-null  float64
2   origin                               1198458 non-null  object
3   dest                                 1198458 non-null  object
4   uniquecarrier                        1198458 non-null  object
5   marketshareorigin                    1198458 non-null  float64
6   marketsharedest                      1198458 non-null  float64
7   hhiorigin                            1198458 non-null  float64
8   hhidest                              1198458 non-null  float64
9   nonhubairportorigin                  1198458 non-null  int64

```

10	smallhubairportorigin	1198458	non-null	int64
11	mediumhubairportorigin	1198458	non-null	int64
12	largehubairportorigin	1198458	non-null	int64
13	nonhubairportdest	1198458	non-null	int64
14	smallhubairportdest	1198458	non-null	int64
15	mediumhubairportdest	1198458	non-null	int64
16	largehubairportdest	1198458	non-null	int64
17	nonhubairlineorigin	1198458	non-null	int64
18	smallhubairlineorigin	1198458	non-null	int64
19	mediumhubairlineorigin	1198458	non-null	int64
20	largehubairlineorigin	1198458	non-null	int64
21	nonhubairlinedest	1198458	non-null	int64
22	smallhubairlinedest	1198458	non-null	int64
23	mediumhubairlinedest	1198458	non-null	int64
24	largehubairlinedest	1198458	non-null	int64
25	year	1198458	non-null	int64
26	month	1198458	non-null	int64
27	dayofmonth	1198458	non-null	int64
28	dayofweek	1198458	non-null	int64
29	scheduledhour	1198458	non-null	int64
30	tailnum	1198458	non-null	object
31	capacity	1198458	non-null	int64
32	loadfactor	1198458	non-null	float64
33	numflights	1198458	non-null	float64
34	distance	1198458	non-null	int64
35	monopolyroute	1198458	non-null	int64
36	temperature	1198458	non-null	float64
37	temp_ninfty_n10	1198458	non-null	int64
38	temp_n10_0	1198458	non-null	int64
39	temp_0_10	1198458	non-null	int64
40	temp_10_20	1198458	non-null	int64
41	temp_20_30	1198458	non-null	int64
42	temp_30_40	1198458	non-null	int64
43	temp_40_infty	1198458	non-null	int64
44	windspeed	1198458	non-null	float64
45	windspeedsquare	1198458	non-null	float64
46	windgustdummy	1198458	non-null	int64
47	windgustspeed	1198458	non-null	float64
48	raindummy	1198458	non-null	int64
49	raintracedummy	1198458	non-null	int64
50	snowdummy	1198458	non-null	int64
51	snowtracedummy	1198458	non-null	int64
52	originmetropop	1198458	non-null	int64
53	originmetrogdppercapita	1198458	non-null	float64
54	destmetropop	1198458	non-null	int64
55	destmetrogdppercapita	1198458	non-null	float64

dtypes: float64(13), int64(39), object(4)

memory usage: 521.2+ MB

5 Rectifying incorrectly encoded ordinal categorical variables

As deduced from the initial analysis, there are few columns that are already one-hot encoded in the dataset. However, some of these categorical variables (described below) are ordinal and one-hot encoding is not appropriate for them.

- The columns of the form `temp_<lower-limit>_<upper-limit>` denote the range in which `temperature` falls in. These columns hold 1 if the `temperature` falls in that range and 0 otherwise. However, the ranges have an order associated with them as follows:

$-\infty$ to $-10 < -10$ to $0 < 0$ to $10 < 10$ to $20 < 20$ to $30 < 30$ to $40 < 40$ to ∞

Hence, a single column `temp_range` is created and contains the categories. All the columns already present as a part of one-hot encoding are dropped.

- The columns of the form `<size of hub>airportorigin` and `<size of hub>airportdest` denote whether the origin and destination airports are hubs for some airline and if its a hub what is its size. These columns hold 1 if the `origin` or `dest` is a `<size of hub>` hub for some airline and 0 otherwise.

`nonhub < small < medium < large`

Hence, a two columns `hubairportorigin` and `hubairportdest` are created and contains the categories. All the columns already present as a part of one-hot encoding are dropped.

- The columns of the form `<size of hub>airlineorigin` and `<size of hub>airlinedest` denote whether the origin and destination airports are hubs for the airline and if its a hub what is its size. These columns hold 1 if the `origin` or `dest` is a `<size of hub>` hub for `uniquecarrier` and 0 otherwise.

`nonhub < small < medium < large`

Hence, a two columns `hubairlineorigin` and `hubairlinedest` are created and contains the categories. All the columns already present as a part of one-hot encoding are dropped.

```
[12]: # extracting all dummy features from the respective categories.
```

```
cols = list(flight_df.columns)
temperature_range = cols[37:44]
airport_connectivity_origin = cols[9:13]
airport_connectivity_dest = cols[13:17]
airline_connectivity_origin = cols[17:21]
airline_connectivity_dest = cols[21:25]
```

```
[13]: # converts dummy columns into a single categorical feature.
```

```
def onehot2ordinal(new_colname: str, dummies: list, str2replace: str, sep = "\n",
                  ↪None):

    flight_df[new_colname] = pd.from_dummies(flight_df[dummies],
                                              default_category = np.nan,
                                              sep = sep)

    if not sep:
```

```

        flight_df[new_colname] = flight_df[new_colname].astype(str).apply(lambda x:
↪x: x.replace(str2replace,

        ''))

flight_df.drop(dummies, axis = 1, inplace = True)

```

```

[14]: # applies the function to all the respective kinds of dummy features.
arguments = [{"temp_range", temperature_range, None, "_"},
↪["hubairportorigin", airport_connectivity_origin, "airportorigin",
↪None],
["hubairportdest", airport_connectivity_dest, "airportdest", None],
["hubairlineorigin", airline_connectivity_origin, "airlineorigin",
↪None],
["hubairlinedest", airline_connectivity_dest, "airlinedest", None]]

for new_colname, dummies, str2replace, sep in arguments:

    onehot2ordinal(new_colname = new_colname, dummies = dummies,
                    str2replace = str2replace, sep = sep)

```

It appears that `temp_range` has some null values and these can be derived by categorizing the temperature value for the particular row.

```

[15]: def categorise_temp(temp):
        if temp <= -10:
            return "ninfty_n10"
        elif temp > -10 and temp <= 0:
            return "n10_0"
        elif temp > 0 and temp <= 10:
            return "0_10"
        elif temp > 10 and temp <= 20:
            return "10_20"
        elif temp > 20 and temp <= 30:
            return "20_30"
        elif temp > 30 and temp <= 40:
            return "30_40"
        else:
            return "40_infty"

```

```

[16]: flight_df["temp_range"] = flight_df.apply(lambda x: categorise_temp(x.
↪temperature),

                                                axis = 1)

```

```

[17]: flight_df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Index: 1198458 entries, 0 to 1201663

```


Data columns (total 38 columns):

#	Column	Non-Null Count	Dtype
0	depdelay	1198458 non-null	int64
1	arrdelay	1198458 non-null	float64
2	origin	1198458 non-null	object
3	dest	1198458 non-null	object
4	uniquecarrier	1198458 non-null	object
5	marketshareorigin	1198458 non-null	float64
6	marketsharedest	1198458 non-null	float64
7	hhiorigin	1198458 non-null	float64
8	hhidest	1198458 non-null	float64
9	year	1198458 non-null	int64
10	month	1198458 non-null	int64
11	dayofmonth	1198458 non-null	int64
12	dayofweek	1198458 non-null	int64
13	scheduledhour	1198458 non-null	int64
14	tailnum	1198458 non-null	object
15	capacity	1198458 non-null	int64
16	loadfactor	1198458 non-null	float64
17	numflights	1198458 non-null	float64
18	distance	1198458 non-null	int64
19	monopolyroute	1198458 non-null	int64
20	temperature	1198458 non-null	float64
21	windspeed	1198458 non-null	float64
22	windspeedsquare	1198458 non-null	float64
23	windgustdummy	1198458 non-null	int64
24	windgustspeed	1198458 non-null	float64
25	raindummy	1198458 non-null	int64
26	raintracedummy	1198458 non-null	int64
27	snowdummy	1198458 non-null	int64
28	snowtracedummy	1198458 non-null	int64
29	originmetropop	1198458 non-null	int64
30	originmetrogdppercapita	1198458 non-null	float64
31	destmetropop	1198458 non-null	int64
32	destmetrogdppercapita	1198458 non-null	float64
33	temp_range	1198458 non-null	object
34	hubairportorigin	1198458 non-null	object
35	hubairportdest	1198458 non-null	object
36	hubairlineorigin	1198458 non-null	object
37	hubairlinedest	1198458 non-null	object

dtypes: float64(13), int64(16), object(9)

memory usage: 356.6+ MB

6 Binning the delays into categories

- The arrival and departure delay are binned into 5 categories.

- The arrival delay is the target variable and will be predicted using classification algorithms.

```
[18]: def bin_delay(delay):
      if delay <= 10:
          return "0-10"
      elif 10 < delay <= 20:
          return "10-20"
      elif 20 < delay <= 40:
          return "20-40"
      elif 40 < delay <= 60:
          return "40-60"
      elif delay > 60:
          return ">60"
```

```
[19]: flight_df["depdelay"] = flight_df.apply(lambda x: bin_delay(x.depdelay),
                                             axis = 1)
      flight_df["arrdelay"] = flight_df.apply(lambda x: bin_delay(x.arrdelay),
                                             axis = 1)
```

7 Splitting the data into train-test

```
[20]: target = "arrdelay"
```

```
[21]: # encoding the target variable
      flight_df[target] = flight_df[target].replace(dict(zip(['0-10', '10-20',
      ↪ '20-40', '40-60', '>60'],
      [i for i in range(5)])))
```

```
[22]: flight_y = flight_df.pop(target)
```

```
[23]: flight_X_train, flight_X_test, flight_y_train, flight_y_test =
      ↪ train_test_split(flight_df, flight_y,
      ↪ test_size = 0.2,
      ↪ stratify = flight_y,
      ↪ random_state = 42)
```

8 Encoding the categorical variables

```
[24]: ord_cat_features = ["temp_range", "hubairportorigin", "hubairportdest",
      "hubairlinedest", "hubairlineorigin", "depdelay"]

      hc_cat_features = ["origin", "dest", "uniquecarrier", "tailnum"]
```

```
cat_vars = ord_cat_features + hc_cat_features
```

8.1 Ordinal variable encoding

```
[25]: categories = [
    ['ninfty_n10', 'n10_0', '0_10', '10_20', '20_30', '30_40', '40_infty'],
    ['nonhub', 'smallhub', 'mediumhub', 'largehub'],
    ['nonhub', 'smallhub', 'mediumhub', 'largehub'],
    ['nonhub', 'smallhub', 'mediumhub', 'largehub'],
    ['nonhub', 'smallhub', 'mediumhub', 'largehub'],
    ['0-10', '10-20', '20-40', '40-60', '>60']
]

oe = OrdinalEncoder(categories = categories).
    ↪fit(flight_X_train[ord_cat_features])

flight_X_train[ord_cat_features] = oe.
    ↪transform(flight_X_train[ord_cat_features])
flight_X_train[ord_cat_features] = flight_X_train[ord_cat_features].
    ↪astype("int64")

flight_X_test[ord_cat_features] = oe.transform(flight_X_test[ord_cat_features])
flight_X_test[ord_cat_features] = flight_X_test[ord_cat_features].
    ↪astype("int64")
```

8.2 High Cardinality variable encoding

- The high cardinality variables in this context include origin, dest, tailnum, and uniquecarrier.
- It is so possible that an origin airport appears in the test data but does not appear in the train dataset.
- Therefore, all the categories must be passed as input to the TargetEncoder() to learn that these values are not in the train dataset.

```
[26]: categories = [list(set(flight_df[i]))for i in hc_cat_features]
```

```
[27]: for hc_col_idx in range(len(hc_cat_features)):
    hc_col = hc_cat_features[hc_col_idx]
    te = TargetEncoder(categories = [categories[hc_col_idx]])
    te = te.fit(flight_X_train[[hc_col]],
                y = flight_y_train)

    flight_X_train[hc_col] = te.transform(flight_X_train[[hc_col]])
    flight_X_test[hc_col] = te.transform(flight_X_test[[hc_col]])
```

```
[28]: flight_X_train.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Index: 958766 entries, 898074 to 382376
Data columns (total 37 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   depdelay                             958766 non-null  int64
1   origin                              958766 non-null  float64
2   dest                                958766 non-null  float64
3   uniquecarrier                       958766 non-null  float64
4   marketshareorigin                   958766 non-null  float64
5   marketsharedest                     958766 non-null  float64
6   hhiorigin                           958766 non-null  float64
7   hhidest                             958766 non-null  float64
8   year                                958766 non-null  int64
9   month                               958766 non-null  int64
10  dayofmonth                           958766 non-null  int64
11  dayofweek                            958766 non-null  int64
12  scheduledhour                        958766 non-null  int64
13  tailnum                             958766 non-null  float64
14  capacity                             958766 non-null  int64
15  loadfactor                           958766 non-null  float64
16  numflights                           958766 non-null  float64
17  distance                             958766 non-null  int64
18  monopolyroute                       958766 non-null  int64
19  temperature                          958766 non-null  float64
20  windspeed                           958766 non-null  float64
21  windspeedsquare                     958766 non-null  float64
22  windgustdummy                       958766 non-null  int64
23  windgustspeed                       958766 non-null  float64
24  raindummy                           958766 non-null  int64
25  raintracedummy                     958766 non-null  int64
26  snowdummy                           958766 non-null  int64
27  snowtracedummy                     958766 non-null  int64
28  originmetropop                      958766 non-null  int64
29  originmetrogdppercapita             958766 non-null  float64
30  destmetropop                        958766 non-null  int64
31  destmetrogdppercapita               958766 non-null  float64
32  temp_range                           958766 non-null  int64
33  hubairportorigin                    958766 non-null  int64
34  hubairportdest                      958766 non-null  int64
35  hubairlineorigin                    958766 non-null  int64
36  hubairlinedest                      958766 non-null  int64
dtypes: float64(16), int64(21)
memory usage: 278.0 MB

```

```
[29]: flight_X_train = flight_X_train[flight_X_train.columns]
```

9 Scaling the numerical variables

```
[30]: float64_vars = flight_X_train.select_dtypes(include = ["float64"])
      num_var = list(set(float64_vars).difference(hc_cat_features))
```

```
[31]: num_var
```

```
[31]: ['loadfactor',
      'windspeedsquare',
      'temperature',
      'originmetrogdppercapita',
      'destmetrogdppercapita',
      'marketsharedest',
      'windspeed',
      'marketshareorigin',
      'hhidest',
      'hhiorigin',
      'windgustspeed',
      'numflights']
```


```
[32]: sscler = StandardScaler().fit(flight_X_train[num_var])

      scaled_train = sscler.transform(flight_X_train[num_var])
      flight_X_train[num_var] = scaled_train

      scaled_test = sscler.transform(flight_X_test[num_var])
      flight_X_test[num_var] = scaled_test
```

10 Correlation between the columns

```
[33]: plt.figure(figsize = (33, 25))
      sns.heatmap(flight_X_train.corr(), annot = True)
      plt.xlabel("features", fontsize = 15)
      plt.ylabel("features", fontsize = 15)
      plt.title("Correlation Matrix of all numeric features", fontsize=22)
      plt.show()
```



- The **depdelay** has slight positive correlation of 0.19 with **scheduledhour**. This makes sense because the air traffic at the **scheduledhour** decides whether the ATC gives permission for aircraft to depart.
- Moreover, there is very slight correlation with the columns pertaining to weather. This also makes sense since the weather (visibility) at departure decides whether the plane can take-off or not.
- There is positive correlation of 0.44 between **tailnum** and **uniquecarrier** since a particular aircraft is owned by a particular carrier. Hence, that correlation does make sense.
- All the weather columns have some relation with each other.
- The **temperature** and **temp_range** have a very high positive correlation of 0.96 and it does make sense as well. Therefore, one of the columns needs to be dropped.
- On similar lines, **windspeed** and **windspeedsquare** have very high positive correlation of 0.93 and hence one of them needs to be dropped.
- The **capacity** and **loadfactor** have a positive correlation of approximately 0.40 with the **distance**. This makes sense because long-haul routes generally have higher **capacity** and a higher payload resulting a higher **loadfactor**.
- Not all values in the matrix make sense. They arise as an artifact of the data but do not make sense at all.

```
[34]: flight_X_train = flight_X_train.drop(["windspeedsquare", "temperature"], axis = 1)
      flight_X_test = flight_X_test.drop(["windspeedsquare", "temperature"], axis = 1)
```

- I dropped windspeedsquare since windspeed is more intuitive to understand and explain the models that will be trained later.
- The temperature column is dropped. This is because the target is a range of delay and keep the temp_range might help instead of actual value of temperature.

```
[35]: from sklearn.tree import DecisionTreeClassifier
      from sklearn.model_selection import train_test_split
      from sklearn.metrics import classification_report, accuracy_score
      import pandas as pd

      # Initialize the decision tree
      clf = DecisionTreeClassifier(max_depth=5, random_state=42)
      clf.fit(flight_X_train, flight_y_train)

      # Predict and Evaluate
      y_pred = clf.predict(flight_X_test)
      accuracy = accuracy_score(flight_y_test, y_pred)
      report = classification_report(flight_y_test, y_pred)

      # Print the result
      print("Decision Tree results:")
      print(f"Accuracy: {accuracy}")
      print(f"Classification Report:\n{report}")
```

Decision Tree results:

Accuracy: 0.8510672029104017

Classification Report:

	precision	recall	f1-score	support
0	0.89	0.99	0.94	183773
1	0.43	0.03	0.05	18537
2	0.48	0.47	0.47	16565
3	0.49	0.43	0.46	7494
4	0.88	0.85	0.86	13323
accuracy			0.85	239692
macro avg	0.63	0.55	0.56	239692
weighted avg	0.82	0.85	0.82	239692

```
[36]: from sklearn.ensemble import HistGradientBoostingClassifier
      from sklearn.model_selection import GridSearchCV
      import time
```

```

start_time = time.time()

# hyperparameter tuning
param_grid = {
    "learning_rate": [0.01, 0.1, 0.2],
    "max_iter": [200, 400, 600, 800, 1000],
    "max_depth": [8, 16, 24]
}

# Grid search
hgbc = HistGradientBoostingClassifier(random_state=42)
grid_search = GridSearchCV(hgbc, param_grid=param_grid, cv=5,
    ↪scoring='accuracy', n_jobs=-1)
grid_search.fit(flight_X_train, flight_y_train)

# Print out the best hyperparameter for future training reference
print(f"Best hyperparameters from GridSearchCV:")
for param, value in grid_search.best_params_.items():
    print(f"  {param}: {value}")

best_hgbc = grid_search.best_estimator_
hgbc_test_score = best_hgbc.score(flight_X_test, flight_y_test)
print(f"HistGradientBoostingClassifier test accuracy: {hgbc_test_score}")

# Predict, and Evaluate
y_pred_hgbc = best_hgbc.predict(flight_X_test)
print("HistGradientBoostingClassifier results:")
HGBC_accuracy = accuracy_score(flight_y_test, y_pred_hgbc)
print(f"HGBC Accuracy: {HGBC_accuracy}")

print("HistGradientBoostingClassifier results:")
print(classification_report(flight_y_test, y_pred_hgbc))

# time the entire experiment
end_time = time.time()
time_to_select = end_time - start_time
print("total run time: ");
print(time_to_select)

```

Best hyperparameters from GridSearchCV:

```

learning_rate: 0.01
max_depth: 8
max_iter: 1000

```

HistGradientBoostingClassifier test accuracy: 0.8519141231246766

HistGradientBoostingClassifier results:

HGBC Accuracy: 0.8519141231246766

HistGradientBoostingClassifier results:

	precision	recall	f1-score	support
0	0.89	0.99	0.94	183773
1	0.42	0.02	0.04	18537
2	0.48	0.47	0.48	16565
3	0.51	0.40	0.45	7494
4	0.87	0.86	0.86	13323
accuracy			0.85	239692
macro avg	0.63	0.55	0.55	239692
weighted avg	0.81	0.85	0.82	239692

total run time:
5801.750147104263

[]:

[]: